

# Simple Linear Regression

---

In this example, we will consider sales based on 'TV' marketing budget.  
We'll build a linear regression model to predict 'Sales' using 'TV' as the predictor variable.

## Understanding the data

### Import pandas

```
In [1]: import pandas as pd
```

### Read the dataset with pandas

```
In [2]: df = pd.read_csv('tvmarketing.csv')
```

### View the first five rows of the dataset

```
In [3]: df.head()
```

```
Out[3]:
```

	TV	Sales
0	230.1	22.1
1	44.5	10.4
2	17.2	9.3
3	151.5	18.5
4	180.8	12.9

### View the last 5 rows

```
In [4]: df.tail()
```

```
Out[4]:
```

	TV	Sales
195	38.2	7.6
196	94.2	9.7
197	177.0	12.8
198	283.6	25.5
199	232.1	13.4

### Check info about the dataset

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    TV      200 non-null    float64
1   Sales   200 non-null    float64
dtypes: float64(2)
memory usage: 3.2 KB
```

### Check the number of rows and columns in the dataset

```
In [6]: df.shape
```

```
Out[6]: (200, 2)
```

### Quick view the basic statistical information about the dataset

```
In [7]: df.describe()
```

```
Out[7]:
```

	TV	Sales
count	200.000000	200.000000
mean	147.042500	14.022500
std	85.854236	5.217457
min	0.700000	1.600000
25%	74.375000	10.375000
50%	149.750000	12.900000
75%	218.825000	17.400000
max	296.400000	27.000000


# Visualising the data using Seaborn Library

Import Matplotlib, Seaborn and set matplotlib inline

```
In [8]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Visualize the relationship between the features and the response using scatterplots

```
In [9]: plt.figure(figsize = (4,6))
sns.scatterplot(x = "TV", y = "Sales", data = df)
plt.show()
```



## Performing Simple Linear Regression

Equation of linear regression

$$y = c + m_1x_1 + m_2x_2 + \dots + m_nx_n$$

- $y$  is the response
- $c$  is the intercept
- $m_1$  is the coefficient for the first feature
- $m_n$  is the coefficient for the nth feature

In our case:

$$y = c + m_1 \times TV$$

The  $m_1$  value is called the model coefficient or model parameter.

## Generic Steps in Model Building using sklearn

Before you read further, it is good to understand the generic structure of modeling using the scikit-learn library. Broadly, the steps to build any model can be divided as follows:

### Preparing X and y (Independent and Dependent variables)

- The scikit-learn library expects X (feature variable) and y (response variable) to be NumPy arrays.
- However, X can be a dataframe as Pandas is built over NumPy.

#### Assign feature variable(Independent variable) to X

```
In [10]: X = df[['TV']]
```

#### Print the first five rows of X

```
In [11]: X.head()
```

```
Out[11]:
```

	TV
0	230.1
1	44.5
2	17.2
3	151.5
4	180.8

#### Assign response variable(Dependent variable or target variable) to y

```
In [12]: y = df['Sales']
```

#### Split the data into Training and testing Sets

```
In [13]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y , test_size = 0.3, ran
```

```
In [14]: # random_state is the seed used by the random number generator, it can be any  
# If we don't specify a random state integer, we get different training and te
```

```
In [ ]: train_test_split    #Press Tab to auto-fill the code  
#Press Tab+Shift to read the documentation
```

### Check the shape of training and testing set

```
In [20]: print(X_train.shape)  
print(y_train.shape)  
print(X_test.shape)  
print(y_test.shape)
```

```
(140, 1)  
(140,)  
(60, 1)  
(60,)
```

## Performing Linear Regression

### Import LinearRegression from sklearn linear model

```
In [21]: from sklearn.linear_model import LinearRegression
```

### Create a LinearRegression object

```
In [22]: lr = LinearRegression()
```

### Train the model (lr) using fit

```
In [23]: lr.fit(X_train, y_train)
```

```
Out[23]: LinearRegression()
```

### Print the coefficient and Intercept

```
In [24]: print(lr.coef_)  
print(lr.intercept_)
```

```
[0.04649736]  
6.989665857411679
```

### Check the score of our model

```
In [27]: lr.score(X_train, y_train)
```

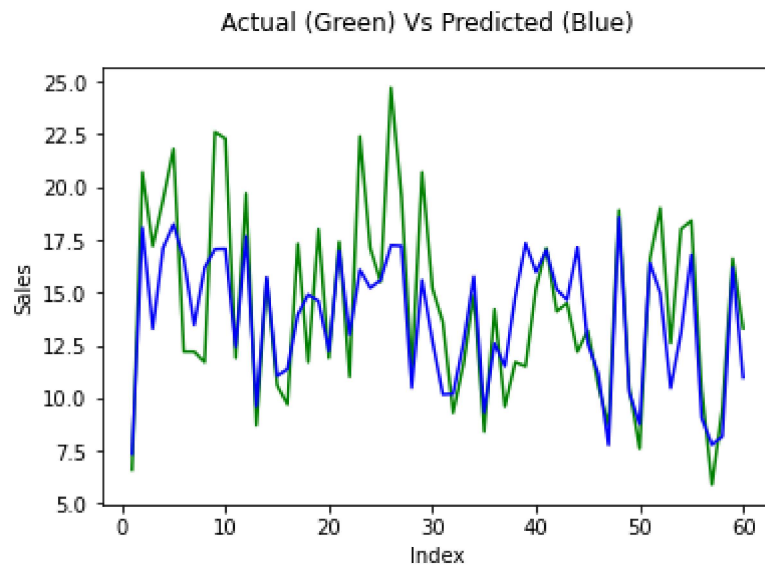
```
Out[27]: 0.6134312469429994
```

## Make predictions on the testing set

```
In [25]: y_pred = lr.predict(X_test)
```


## Plot a graph to check the accuracy of our prediction

```
In [26]: c = [i for i in range(1,61,1)]    #Creating an Index, 61 is used because we hav  
fig = plt.figure()  
plt.plot(c,y_test, color = 'green') # Plotting y test  
plt.plot(c,y_pred, color = 'blue') # Plotting predicted values  
fig.suptitle('Actual (Green) Vs Predicted (Blue)') # Set title  
plt.xlabel('Index') # Set X Label  
plt.ylabel('Sales') # Set Y Label  
plt.show()
```



## Plot a scatterplot of actual values vs predicted

```
In [36]: sns.scatterplot(x = y_test, y = y_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predicted values')
plt.show()
```



## Calculate mean squared error and r2\_score

Mean Squared Error: It is the sum, over all the data points, of the square of the difference between the predicted and actual target variables, divided by the number of data points.

$y - y_{pred}^2$

$$e_i = Y_{Pred} - Y_i$$

$$Y_{pred} = mx + c$$

i.e,

$$e_1 = ((m_i * x + c) - y_i)$$

i.e,

$$e_1^2 = (y_1 - (m_i * x + c))^2$$

Import mean\_squared\_error and r2\_score from sklearn.metrics

```
In [30]: from sklearn.metrics import mean_squared_error, r2_score
```

Calculate the mean squared error

In [33]: `mean_squared_error(y_test, y_pred)`

Out[33]: 7.97579853285485

### Calculate r2\_score

In [34]: `r2_score(y_test, y_pred)`

Out[34]: 0.5942987267783302

---